

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.tree import plot_tree
from sklearn import tree
from sklearn import metrics
from sklearn import svm
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import roc_auc_score
from sklearn.metrics import accuracy_score
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsRegressor
from sklearn.model_selection import train_test_split

import warnings
warnings.filterwarnings('ignore')

df = pd.read_csv('titanic.csv')
df.head()

```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	SibSp	\	Name	Sex	Age
0			Braund, Mr. Owen Harris	male	22.0
1					
1	1		Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0
1					
2			Heikkinen, Miss. Laina	female	26.0
0					
3			Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0
1					
4			Allen, Mr. William Henry	male	35.0
0					

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```
df.size
```

```
10692
```

```
df.shape
```

```
(891, 12)
```

```
df.columns
```

```
Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age',
      'SibSp',
      'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],
      dtype='object')
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 891 entries, 0 to 890
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	PassengerId	891 non-null	int64
1	Survived	891 non-null	int64
2	Pclass	891 non-null	int64
3	Name	891 non-null	object
4	Sex	891 non-null	object
5	Age	714 non-null	float64
6	SibSp	891 non-null	int64
7	Parch	891 non-null	int64
8	Ticket	891 non-null	object
9	Fare	891 non-null	float64
10	Cabin	204 non-null	object
11	Embarked	889 non-null	object

```
dtypes: float64(2), int64(5), object(5)
```

```
memory usage: 83.7+ KB
```

```
df.isnull().sum()
```

PassengerId	0
Survived	0
Pclass	0
Name	0
Sex	0

```
Age          177
SibSp         0
Parch         0
Ticket        0
Fare          0
Cabin        687
Embarked      2
dtype: int64
```

```
# excluding collums which is not required
df.drop(["PassengerId",
"Name","SibSp","Parch","Ticket","Fare","Cabin","Embarked"],
axis="columns",inplace=True)
df.head()
```

	Survived	Pclass	Sex	Age
0	0	3	male	22.0
1	1	1	female	38.0
2	1	3	female	26.0
3	1	1	female	35.0
4	0	3	male	35.0

```
X = df.drop(["Survived"],axis="columns")
Y = df["Survived"]
Y
```

0	0
1	1
2	1
3	1
4	0
..	
886	0
887	1
888	0
889	1
890	0

```
Name: Survived, Length: 891, dtype: int64
```

```
X['Sex'] = X['Sex'].replace({"male": 0, "female": 1})
X.head()
```

	Pclass	Sex	Age
0	3	0	22.0
1	1	1	38.0
2	3	1	26.0
3	1	1	35.0
4	3	0	35.0

```
X.Age[0:10]
```

0	22.0
1	38.0
2	26.0
3	35.0
4	35.0
5	NaN
6	54.0
7	2.0
8	27.0
9	14.0

Name: Age, dtype: float64

```
X['Age'].fillna(X['Age'].mean(), inplace=True)
```

```
X['Age'].iloc[0:10]
```

0	22.000000
1	38.000000
2	26.000000
3	35.000000
4	35.000000
5	29.699118
6	54.000000
7	2.000000
8	27.000000
9	14.000000

Name: Age, dtype: float64

```
train_test_split(X,Y,train_size = 0.8)
```

	Pclass	Sex	Age
629	3	0	29.699118
721	3	0	17.000000
251	3	1	29.000000
105	3	0	28.000000
693	3	0	25.000000
..	...	...	...
471	3	0	38.000000
476	2	0	34.000000
691	3	1	4.000000
767	3	1	30.500000
11	1	1	58.000000

[712 rows x 3 columns],

	Pclass	Sex	Age
764	3	0	16.000000
121	3	0	29.699118
630	1	0	80.000000
393	1	1	23.000000
182	3	0	9.000000
..	...	...	...

841	2	0	16.000000
36	3	0	29.699118
311	1	1	18.000000
572	1	0	36.000000
277	2	0	29.699118

[179 rows x 3 columns],

629	0
721	0
251	0
105	0
693	0

..	
471	0
476	0
691	1
767	0
11	1

Name: Survived, Length: 712, dtype: int64,

764	0
121	0
630	1
393	1
182	0

..	
841	0
36	1
311	1
572	1
277	0

Name: Survived, Length: 179, dtype: int64]

```
X_train,X_test,Y_train,Y_real=train_test_split(X,Y,train_size = 0.8)
len(X_train)
```

712

```
len(X_test)
```

179

X\_test

	Pclass	Sex	Age
250	3	0	29.699118
507	1	0	29.699118
212	3	0	22.000000
344	2	0	36.000000
3	1	1	35.000000
..	...	...	...
440	2	1	45.000000

38	3	1	18.000000
113	3	1	20.000000
237	2	1	8.000000
812	2	0	35.000000

[179 rows x 3 columns]

# Model Building

## 1. Decision Tree

```

model_dt = tree.DecisionTreeClassifier(max_depth=3)
model_dt

DecisionTreeClassifier(max_depth=3)

model_dt = DecisionTreeClassifier()
model_dt.fit(X_train, Y_train)
predictions = model_dt.predict(X_test)

model_dt.predict(X_test)

array([0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0,
0,
      0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1,
0,
      1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0,
1,
      0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1,
0,
      0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0,
0,
      0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
0,
      0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0,
1,
      0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1,
0,
      0, 1, 0], dtype=int64)

Y_pred_dt = model_dt.predict(X_test)
Y_pred_dt

array([0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0,
0,
      0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1,
0,
      0,

```

```

1,      1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0,
0,      0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1,
0,      0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0,
0,      0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
1,      0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0,
0,      0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1,
0,      0, 1, 0], dtype=int64)

```

```

result = pd.DataFrame({"Survived REAL": Y_real, "Survived PREDICTION":
predictions})
result

```

	Survived REAL	Survived PREDICTION
250	0	0
507	1	0
212	0	0
344	0	0
3	1	1
..	...	...
440	1	1
38	0	0
113	0	0
237	1	1
812	0	0

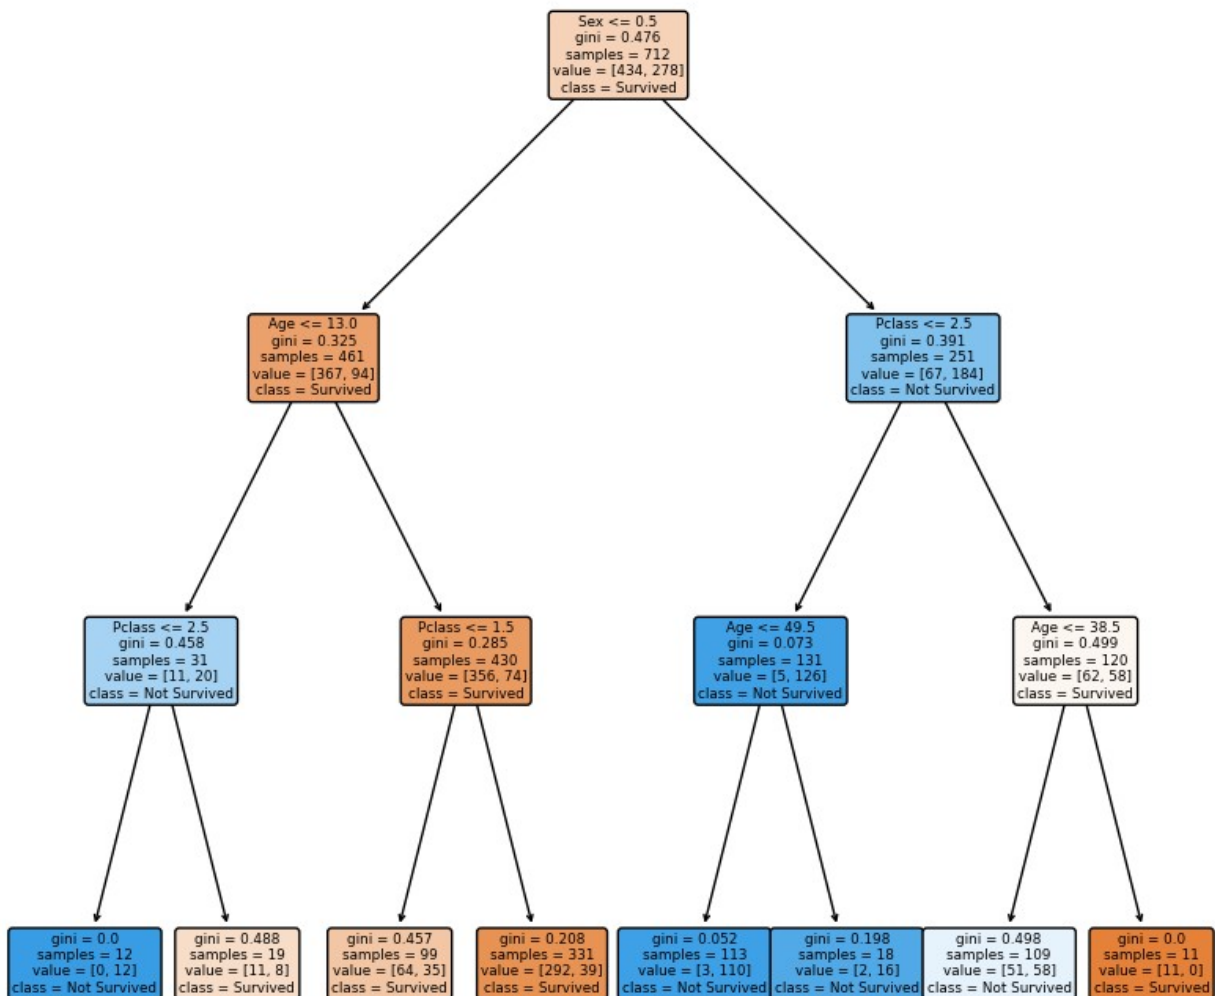
```
[179 rows x 2 columns]
```

## Plotting Tree

```

model_dt = DecisionTreeClassifier(max_depth=3)
model_dt.fit(X_train, Y_train)
plt.figure(figsize=(10, 10))
plot_tree(model_dt, feature_names=['Pclass', 'Sex', 'Age'],
class_names=['Survived', 'Not Survived'], filled=True, rounded=True)
plt.show()

```



## Performance Measurement

### Accuracy

```

accuracy_dt = model_dt.score(X_test, Y_real)
accuracy_decimal = "{:.3f}".format(accuracy_dt)
print("Accuracy:", accuracy_decimal)

```

Accuracy: 0.832



## Precision

```
precision_dt = metrics.precision_score(Y_real, Y_pred_dt)
precision_decimal = "{:.3f}".format(precision_dt)
print("Precision:", precision_decimal)
```

Precision: 0.700

## Recall

```
recall_dt = metrics.recall_score(Y_real, Y_pred_dt)
recall_decimal = "{:.3f}".format(recall_dt)
print("Recall:", recall_decimal)
```

Recall: 0.656


## Confusion Matrix

```
confusion_matrix_dt = metrics.confusion_matrix(Y_real, Y_pred_dt)
confusion_matrix_dt
```

```
array([[97, 18],
       [22, 42]], dtype=int64)
```

```
sns.heatmap(confusion_matrix_dt, annot=True)
```

<Axes: >



## AUC - ROC Curve

```
auc_dt = metrics.roc_auc_score(Y_real, Y_pred_dt)
auc_dt_formatted = f'{auc_dt:.3f}'
print("AUC Score:", auc_dt_formatted)
```

AUC Score: 0.750

## DT K-Fold Cross Validation

```
kf_dt = KFold(n_splits=5, shuffle=True)
cv_dt = cross_val_score(model, X, Y, cv=kf_dt)
mean_cv_score = np.mean(cv_dt)
mean_cv_score_formatted = f'{mean_cv_score:.3f}'
print("Mean Cross-Validation Score:", mean_cv_score_formatted)
```

Mean Cross-Validation Score: 0.818

## Logistic Regression

```
model_lr = LogisticRegression(solver='liblinear')
```

```

model_lr.fit(X_train,Y_train)
LogisticRegression(solver='liblinear')
y_pred_lg = model_lr.predict(X_test)
y_pred_lg
array([0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0,
0,
      0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1,
0,
      1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0,
1,
      0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1,
0,
      1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0,
0,
      0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1,
      0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 0,
0,
      0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1,
1,
      1, 1, 0], dtype=int64)

```

## Performance Measurement

### Accuracy

```

accuracy_lg = model_lr.score(X_test, Y_real)
accuracy_lg_formatted = f'{accuracy_lg:.3f}'
print("Accuracy:", accuracy_lg_formatted)

```

Accuracy: 0.838

### Precision

```

precision_lg = metrics.precision_score(Y_real, y_pred_lg)
precision_lg_percentage = "{:.3f}%".format(precision_lg * 100)
print("Precision:", precision_lg_percentage)

```

Precision: 76.923%

### Recall

```


rec_lg = metrics.recall_score(Y_real, y_pred_lg)
rec_lg_percentage = "{:.3f}%".format(rec_lg * 100)
print("Recall:", rec_lg_percentage)

```

Recall: 78.125%

## Confusion Matrix

```
cm_lg = metrics.confusion_matrix(Y_real,y_pred_lg)
cm_lg
array([[100, 15],
       [ 14, 50]], dtype=int64)
sns.heatmap(cm_lg, annot=True)
<Axes: >
```



## AUC - ROC Curve

```
auc_lg = metrics.roc_auc_score(Y_real, y_pred_lg)
auc_lg_percentage = "{:.3f}%".format(auc_lg * 100)
print("AUC Score:", auc_lg_percentage)
AUC Score: 82.541%
```

## LR K-Fold Cross Validation

```
kf_lg = KFold(n_splits=5,shuffle=True)
```

```

cv_lg = cross_val_score(model_lr, X, Y, cv=kf_lg)
mean_cv_accuracy = np.mean(cv_lg)
mean_cv_accuracy_percentage = "{:.3f}%".format(mean_cv_accuracy * 100)
print("Cross Validation Accuracy:", mean_cv_accuracy_percentage)

```

Cross Validation Accuracy: 79.682%

## Support Vector Machines (SVM)

```

model_svc = svm.SVC(kernel="linear")
model_svc

SVC(kernel='linear')

model_svc.fit(X_train,Y_train)

SVC(kernel='linear')

Y_pred_svm = model_svc.predict(X_test)
Y_pred_svm

array([0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0,
0,
      0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1,
0,
      1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0,
1,
      0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1,
0,
      1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0,
0,
      0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1,
      0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 0,
0,
      0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1,
1,
      1, 1, 0], dtype=int64)

```

## Performance Measurement

### Accuracy

```

accuracy_svc = model_svc.score(X_test, Y_real)
accuracy_svc_formatted = f'{accuracy_svc:.3f}'
print("Accuracy:", accuracy_svc_formatted)

```

Accuracy: 0.838

## Precision

```
precision_svc = precision_score(Y_real, Y_pred_svm)
precision_svc_percentage = "{:.3f}%".format(precision_svc * 100)
print("Precision:", precision_svc_percentage)
```

Precision: 77.778%

## Recall

```
rec_svc = recall_score(Y_real, Y_pred_svm)
rec_svc_percentage = "{:.3f}%".format(rec_svc * 100)
print("Recall:", rec_svc_percentage)
```

Recall: 76.562%


## Confusion Matrix

```
cm_svc = metrics.confusion_matrix(Y_real, Y_pred_svm)
cm_svc
```

```
array([[101, 14],
       [ 15, 49]], dtype=int64)
```

```
sns.heatmap(cm_svc, annot=True)
```

```
<Axes: >
```



## AUC - ROC Curve

```
auc_svc = roc_auc_score(Y_real, Y_pred_svm)
auc_svc_percentage = "{:.3f}%".format(auc_svc * 100)
print("AUC:", auc_svc_percentage)
```

AUC: 82.194%

## SVM K-Fold Cross Validation

```
kf_svc = KFold(n_splits=5, shuffle=True)

cv_svc = cross_val_score(model_svc, X, Y, cv=kf_svc)
mean_cv_svc = np.mean(cv_svc)
mean_cv_svc_percentage = "{:.3f}%".format(mean_cv_svc * 100)
print("Cross-validated Accuracy:", mean_cv_svc_percentage)
```

Cross-validated Accuracy: 78.678%

## Random Forest

```
model_rf = RandomForestClassifier().fit(X_train, Y_train)
```

```

y_pred_forest = model_rf.predict(X_test)
y_pred_forest

array([0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0,
0,
      0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1,
0,
      1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0,
1,
      0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1,
0,
      0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0,
0,
      0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
1,
      0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0,
1,
      0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1,
0,
      0, 1, 0], dtype=int64)

```

## Performance Measurement

### Accuracy

```

accuracy_forest = model_rf.score(X_test, Y_real)
accuracy_forest_formatted = f'{accuracy_forest:.3f}'
print("Accuracy:", accuracy_forest_formatted)

```

Accuracy: 0.793

### Precision

```

precision_forest = precision_score(Y_real, y_pred_forest)
precision_forest_percentage = "{:.3f}%".format(precision_forest * 100)
print("Precision:", precision_forest_percentage)

```

Precision: 71.429%

### Recall

```

recall_forest = recall_score(Y_real, y_pred_forest)
recall_forest_percentage = "{:.3f}%".format(recall_forest * 100)
print("Recall:", recall_forest_percentage)


```

Recall: 70.312%



## Confusion Matrix

```
cm_forest = metrics.confusion_matrix(Y_real, y_pred_forest)
cm_forest
array([[97, 18],
       [19, 45]], dtype=int64)
sns.heatmap(cm_forest, annot=True)
<Axes: >
```



## AUC - ROC Curve

```
auc_forest = roc_auc_score(Y_real, y_pred_forest)
auc_forest_percentage = "{:.3f}%".format(auc_forest * 100)
print("AUC:", auc_forest_percentage)
AUC: 77.330%
```

## RF K-Fold Cross-Validation

```
kf_forest = KFold(n_splits=5, shuffle=True)
cv_forest = cross_val_score(model_rf, X, Y, cv=kf_forest)
mean_accuracy_forest = np.mean(cv_forest)
```

```

mean_accuracy_forest_percentage = "{:.3f}
%".format(mean_accuracy_forest * 100)
print("Cross Validation Accuracy:", mean_accuracy_forest_percentage)

Cross Validation Accuracy: 80.358%

```

## K-Nearest Neighbors

```

knn_model = KNeighborsRegressor()
knn_model

KNeighborsRegressor()

knn_model.fit(X_train, Y_train)

KNeighborsRegressor()

y_pred_knn = knn_model.predict(X_test)
y_pred_knn

array([0. , 0. , 0.2, 0.2, 1. , 0. , 0.8, 0. , 0.2, 0.4, 0.2, 1. ,
0.4,
      1. , 0.2, 1. , 1. , 1. , 0. , 0.2, 0.4, 0. , 0. , 0.2, 0.4,
0.6,
      0. , 0.4, 0.4, 0. , 1. , 1. , 0. , 1. , 0.6, 0. , 0.8, 0.4,
0.6,
      0.2, 0. , 0. , 0.4, 0.4, 0.6, 0.2, 0. , 0.4, 0.4, 0. , 0. ,
0.6,
      0.6, 0.4, 0.6, 0. , 1. , 1. , 1. , 0.6, 0.2, 0.2, 1. , 0.2,
0.2,
      0.4, 0.4, 0. , 0. , 0.2, 0.6, 0.2, 0.6, 0.8, 0.6, 0.2, 0. ,
0. ,
      0.6, 0.2, 0.6, 0.4, 0. , 0.2, 0.4, 0. , 0.4, 0.4, 0. , 0. ,
0. ,
      0.6, 0.4, 0. , 0. , 0. , 0.8, 0. , 0.8, 0.2, 0.8, 0.2, 0.8,
0. ,
      0.2, 0.6, 0.4, 0.8, 0.2, 0. , 0.2, 1. , 1. , 0.4, 0.8, 0.4,
0.2,
      0. , 0.8, 0. , 1. , 0. , 0.2, 0. , 0. , 0. , 0.2, 0. , 0. ,
0.8,
      0.2, 0.6, 0.2, 0.2, 0.2, 0.6, 0. , 0.4, 1. , 0.6, 0.6, 0. ,
0.6,
      0.4, 0.4, 0.4, 0.2, 0.8, 0. , 0.6, 0.6, 1. , 0.4, 0.6, 0.2,
0.2,
      0.4, 1. , 0. , 0.2, 1. , 0.2, 0.6, 0.4, 0.4, 0.2, 0. , 0.2,
0.4,
      0. , 0.6, 0. , 0. , 0.2, 0.4, 0.4, 0.8, 0.4, 0.2])

```

# Performance Measurement

## Accuracy

```
rounding = np.round(y_pred_knn)

acc_knn = metrics.accuracy_score(Y_real, rounding)
acc_knn_formatted = f'{acc_knn:.3f}'
print("Accuracy:", acc_knn_formatted)
```

Accuracy: 0.765

## Precision

```
precision_knn_percentage = precision_score(Y_real, rounding) * 100
print("Precision:", "{:.3f}%".format(precision_knn_percentage))
```

Precision: 68.966%

## Recall

```
recall_knn_percentage = recall_score(Y_real, rounding) * 100
print("Recall:", "{:.3f}%".format(recall_knn_percentage))
```

Recall: 62.500%


## Confusion Matrix

```
cm_knn = metrics.confusion_matrix(Y_real, rounding)
cm_knn

array([[97, 18],
       [24, 40]], dtype=int64)

sns.heatmap(cm_forest, annot=True)

<Axes: >
```



## AUC - ROC Curve

```
auc_knn_percentage = roc_auc_score(Y_real, rounding) * 100  
print("AUC:", "{:.3f}%".format(auc_knn_percentage))
```

AUC: 73.424%

## KNN K-Fold Cross Validation

```
kf_neighbors = KFold(n_splits=5,shuffle=True)  
  
cv_accuracy_neighbors = cross_val_score(knn_model, X, Y,  
cv=kf_neighbors)  
mean_cv_accuracy_neighbors = np.mean(cv_accuracy_neighbors) * 100  
print("Cross Validation Accuracy:", "{:.3f}%".format(mean_cv_accuracy_neighbors))
```

Cross Validation Accuracy: 23.822%

# Valuation Analysis

```
plt.title("Accuracy Comparison Graph")
plt.ylabel("Accuracy Score")

plt.xlabel("Algorithms - 1.Decision Tree / 2.Logistic Regression /
3.SVM (SVC) / 4.Random Forest / 5.KNN")

y = [acc_knn,accuracy_dt,accuracy_lg,accuracy_svc,accuracy_forest]

plt.plot([1,2,3,4,5],y, color = "black")

plt.scatter(1,acc_knn, marker="o", color = "purple",label = "KNN")
plt.scatter(2,accuracy_dt, marker="o", color="red", label="Decision
Tree")
plt.scatter(3,accuracy_lg, marker="o", color="green", label="Logistic
Regression")
plt.scatter(4,accuracy_svc, marker="o", color="blue", label="SVM
(SVC)")
plt.scatter(5,accuracy_forest, marker="o", color = "orange",label =
"Random Forest")

plt.legend()
plt.show()
```

